
MASTER

Research Synthesis

Thirteen pillars. Why this is a portfolio business, not an optimization business.

Respect the Funk · marketing engine research program

Target: ***Losing Sleep*** — Hallow Youth · 15 July 2026

Distils ~6,900 lines and ~2,440 sourced claims across thirteen pillars

The one-paragraph version

Every path to making one song win is closed, and the closures are sourced, cross-confirmed, and mostly not close calls. You cannot buy 100k streams profitably (~50x loss). You cannot measure whether ads caused streams (Facebook couldn't, with 663 RCTs). You cannot advertise film-clip creative (it's a written rule). Bought engagement is EV-negative. Probe accounts need 3–6 months against a rolling 30-day window. A 12-position audience clock costs ~\$118K to validate. Creator demographic data is unaccuracy-rated estimation. And UGC breakout — the last mechanism with non-negative economics — is a **sub-1% lottery** where the top 10% of songs take 96% of creations. **The conclusion those findings force: this is a portfolio business, not an optimization business.** When outcomes are lottery-shaped, the winning move is not a better ticket — it's more tickets at a lower cost per ticket, with the plumbing correct so a win actually pays. *More releases beat more spend.* That is a real engine, and it's buildable: infrastructure is ~\$0 at idle, you own your masters, the highest-value levers are free, and a creative test costs \$143.

1. The strategic verdict

Two strategies were proposed. The evidence supports neither.

	THE PROPOSAL	WHY THE EVIDENCE REJECTS IT
Phase 1 brief	A global, scalable, multi-tenant publishing engine	Scale doesn't fix lottery odds. It multiplies the cost of tickets you're buying at the wrong price.
Phase 2 pivot	Small roster, quality and intelligent spend over quantity, deep audience discovery on one artist	Concentration is the wrong response to a heavy tail. Optimizing one ticket in a sub-1% lottery is the single worst allocation available.

What the evidence actually supports

Many quality shots on goal, with cheap, correct, repeatable plumbing. Not many accounts. Many songs.

This is not a compromise between the two. It's a different axis. The pivot was right that *scale of accounts* is not the lever. It was wrong that *concentration on one song* is — because the payoff distribution makes concentration irrational.

The distribution is the whole argument. [Salganik's Science experiment](#) (14,341 participants, parallel worlds) showed that social influence — precisely what TikTok's feed maximizes — makes outcomes **more unpredictable, not less**. Unpredictability here is not a measurement failure to be engineered away. It is the mechanism. You cannot forecast your way into the tail; you can only buy more draws from it.

What that means for the product

An engine that promises to **make a given song win** is selling something the evidence says nobody can deliver — including the majors, including us.

An engine that makes **each release cheap, correctly-plumbed, fast, and legally clean** is selling something real, and the research says it's buildable for roughly the price of a subscription.

The honest pitch to a label is not "*we make your spend efficient.*" It's "*we make your shots cheap and make sure a win actually pays.*"

2. What is dead, and why

Listed with the evidence, because each of these would have consumed real money before failing.

DEAD	THE FINDING THAT KILLS IT	PILLAR
Buying 100k streams	~\$20–25K (Meta) or ~\$17–79K (TikTok) to produce **~\$300–500** of royalties. A **~50x loss** at any geo or bid strategy.	02, 04, 08
Measuring incrementality	Gordon et al., 663 RCTs inside Facebook , full internal data → <i>"unable to reliably estimate an ad campaign's causal effect."</i>	03
Any verified stream event	Verified by direct fetch of Spotify's own API docs. The complete CAPI event list is VIEW, PRODUCT, CHECK_OUT, ADD_TO_CART, PURCHASE, LEAD, SIGN_UP, CUSTOM_EVENT_1–5 . No stream event exists.	09
Film-clip creative	TikTok's ad policy contains a written rule barring <i>"clips from any unauthorized media sources to promote your product."</i> Meta pre-screens and bans infringing IP. <i>Warhol v. Goldsmith</i> narrows fair use exactly where use is commercial.	06
Bought engagement	EV on a \$500 test: ~\$3,500–14,500 expected cost, ~zero benefit. Cascade term dominates at any probability $\geq 2\%$. Vendors concede buyers <i>"lose 60–80% of inventory within 14 days."</i>	05
Custom ML, year one	<i>"You cannot out-model Meta using a strict subset of Meta's data."</i>	09
The 12-position clock	$C(12,2) = 66$ pairwise tests → \$52,206 naive / ~\$118,500 Bonferroni-corrected. 6–24x the largest budget tier. Mature circumplex literatures top out at 6–8 positions, never 12.	12
Buying creator demographics	No endpoint on TikTok, Meta, or YouTube returns an arbitrary creator's audience demographics at any price. All gated behind that creator's OAuth. Zero of eight vendors publish a demographic accuracy figure.	10
Seeding as a flywheel	No rigorous evidence paid seeding causes organic pickup. Practitioners report momentum ends when spend ends; saturation can suppress adoption.	13

The finding that reframes everything

UGC breakout is a lottery, not a mechanism.

- Only **4%** of 600+ TikTok Top-50 songs crossed to the Hot 100.
- Of ~5,000 TikTok-popular songs, just **125 were emerging artists.**
- Strongest available study (quasi-experiment on UMG's TikTok withdrawal, [arXiv 2405.14999](#)): **top 10% of songs = 96% of creations**, and TikTok's effect **concentrates on already-viral songs — the long tail barely moves.**
- Chained conditionals put a cold indie single **well under 1%.**

The agent commissioned to find the mechanism reported that its own pillar's premise fails. That is the finding working correctly.

An industry that would loudly publish a good base rate publishing none is itself evidence.

3. What survives — the actual engine

After nine sections of "no," this is the part that's real.

The free levers (take all of these first)

LEVER	WHY	COST
TikTok clip start-time	You choose which 60s window plays. The default is wrong. Instagram has no equivalent.	\$0
Metadata hygiene	Pex documents ~20x Sound ID fragmentation . A track that goes viral under fragmented IDs earns nothing .	\$0
You own the masters	RTF owns <i>Losing Sleep</i> . Safe everywhere — organic and paid, all platforms. The business-account music restriction governs <i>other people's</i> music, not yours.	\$0
Be legibly linked	Don't evade Business Manager linkage — the linkage is the alibi , not the risk.	\$0
File the TikTok API audit	No published SLA. The clock only runs once filed.	\$0
Meta needs no App Review	Publishing to accounts your own business owns requires none.	\$0

The cheap instruments

INSTRUMENT	COST	WHAT IT BUYS
One creative test (+50% lift)	\$143	The best-value measurement in the entire stack.
One axis test: People↔Things	\$143–791	Your circle, as a line. See §4.
Sound-usage data (Soundcharts \$10–49/mo, Chartmetric \$40/mo, TikTok sound pages free)	<\$50/mo	Measured behavior instead of estimated demographics.
Creator analytics screenshot	\$0	The same first-party data CreatorIQ charges \$25–60k/yr for. Just ask.
Infrastructure	\$0/mo idle, \$55–366/mo running	A rounding error against any ad spend.
Curator placement	\$50–300/post	Demographic signal in days , not the 3–6 months owned accounts need.

Creative quality is literally a media discount

Meta's own auction docs: $\text{Bid Per Impression} = \text{Bid Per Action} \times \text{Estimated Action Rate}$, and "*an ad that's more relevant could win against ads with higher bids*." Better creative doesn't just convert better — **it buys cheaper impressions**. Meta's own worked example shows pruning an "expensive" placement **raised** blended cost \$3.00 → \$3.25. The manual-optimization instinct is counterproductive per the platform's own documentation.

The whole optimization thesis, honestly stated

*Make more good creative. Feed the machine clean signal. Test only large differences. **Don't fiddle.** Then: **more releases beat more spend.***

4. The circle — the version that survives

The instinct was right and better-founded than it looked. "Math page for men, baby page for women" is almost exactly [Prediger's People](#) [Things dimension](#) — a validated axis underlying Holland's RIASEC hexagon, and one of the most robust findings in vocational psychology. It was reinvented from first principles.

Three reasons it must start smaller than 12 positions:

1. **Cost.** 66 pairwise comparisons = ~\$118,500 corrected. Off by an order of magnitude, and now we know exactly by how much.
2. **Instrument noise.** The demographic data you'd chart with is estimation with **no published accuracy figure from any vendor**. Resolving 12 sectors with an instrument that may not reliably separate 2 is precision theater.
3. **Effect sizes.** Personality [genre](#) correlations run $|r| \approx 0.10-0.21$ — real, but 1–4% of variance. Too weak to route content decisions.

Also: **music-preference research doesn't give a circle**. Rentfrow & Gosling's MUSIC model is 5 factors, and **nobody has run the circularity test on it**. The topology isn't just unfitted at our N — it's unestablished in this domain.

The v1 that works

- **One axis: People** [Things](#). Test as a single pairwise comparison. \$143–791.
- **Earn axis two** (Mellow [Intense](#)) only on signal: \$1,141–6,312 for 4 quadrants.
- **A hexagon is a plausible multi-artist-portfolio goal**. Twelve positions: parked.

Your clock starts as a line.

Two corrections to the design

Don't space probes equidistantly. Real audience density is never uniform (von Mises $\kappa > 0$ in every cited study). **Cluster near the plausible dominant mode** — even spacing wastes probes on empty sectors.

Chart behavior, not demographics. If a creator used an adjacent artist's sound and beat their own median, that's **measured behavioral evidence** their audience likes this music — strictly better than a low-confidence demographic guess, and ~50x cheaper. Same geometry, an instrument that actually measures.

Why the model is worth building at all

Meta's GEM/Andromeda/Lattice stack out-targets anything we could build, on data we'll never have. **So the circle is not a targeting tool**. Its value is an **interpretable vocabulary for what content to make and which creators to seed** — the decision layer the black box doesn't touch, and the one question no ad platform will ever answer.

5. Probe accounts — what the numbers did to the plan

The legitimacy question resolved cleanly in your favor. Meta ships a help article titled "*Managing Multiple Facebook Pages and Instagram Accounts*"; App Review is "**Not required**" for owned accounts. TikTok: "*You can have multiple accounts... but not to deceive others or break the rules.*" **A purpose test, not a headcount test.** Every pillar independently found the same line: **the platforms police deception, not automation.**

The FTC wrinkle is useful: a label speaking as itself isn't a third-party "endorsement" under 16 CFR 255 at all. The real exposure is the **2024 Fake Reviews Rule §465.6** — misrepresenting that an entity you control provides independent opinions.

The timeline question did not resolve in your favor.

FACT	CONSEQUENCE
TikTok 1–5K-follower accounts average 350 views/post, down 59% YoY	The engine barely turns over
Instagram organic reach 3.5%	Same
TikTok Custom Audience needs 1,000 engagers in a rolling 30-day window	Those numbers barely touch
TikTok organic engagement retention is 7/14/30 days ONLY (verified) — Meta's IG accumulates to 730 days	A TikTok probe pool evaporates monthly . It must sustain ~1,000 engagers <i>every month, forever</i> .

Realistic: 3–6 months to signal, 4–8 for a cross-account demographic read.

And the core premise is **UNVERIFIED**: nobody publishes whether a small-seed Lookalike beats Advantage+ broad. You'd spend 3–6 months earning the right to run that test. It is empirical, not researchable, and must not be presented as settled.

The separation that matters

You don't need to own an account to test a demographic — you need access to its audience.

The probe accounts' unique value is the **durable lookalike seed** (and that mostly survives on Instagram, not TikTok). The *discovery signal* is purchasable immediately via curator placement at **\$50–300/post** — days instead of quarters, at the price of a single creative test. **Those are separable, and should be separated.**

Recommendation: 2 accounts, not 8. 90-day gate at 1,000 rolling engagers. On failure, buy curator placement.

The cautionary tale inverts

Chaotic Good drew no ban and no FTC action. But analyst Rachel Karten found their accounts had "**poor engagement metrics**." The risk of fake fan networks isn't that they get punished. **It's that they don't work.**

6. Compliance as product, not overhead

Per [NPR, 13 July 2026](#) — **two days before this research** — labels routinely instruct creators **not to disclose** paid promotion, and were non-responsive to the investigation.

So "do what the majors do" is not available as a compliance strategy. The majors are non-compliant, and it's actively being reported. FTC penalty: \$53,088 per violation.

Build FTC disclosure into creator contracts from day one. This is the rare case where the honest build is also the commercial one: an engine that enforces disclosure by default is a genuine differentiator against an industry that is quietly non-compliant and now under a live investigation.

7. Do this week — ordered, all cheap

1. **Fix the TikTok clip start-time.** Free. The default is wrong. Highest value-per-effort item in the project.
 2. **Audit metadata / Sound ID hygiene.** Free. ~20x fragmentation risk means **a viral track can earn nothing.**
 3. **File the TikTok Content Posting API audit.** Free. No SLA — the clock only starts when filed.
 4. **Ask your distributor about CML clearance.** Free. **Correction: CML does *not* gate organic UGC** — my earlier framing overstated this. It still matters for business-account and ads use.
 5. **Run one People[?]Things test — \$143.** Your circle, as a line. Cheapest real information available.
 6. **Read the Lookalike seed minimums off the live Ads Manager UI.** 60 seconds, and better than any number in these reports (Meta's help pages are JS-only and unfetchable — see §9).
 7. **Decide: portfolio or concentration.** §1 says only one survives. Everything downstream depends on it.
-

8. What we still don't know

GAP	WHY IT MATTERS
Does a small-seed Lookalike beat Advantage+ broad?	The core premise of the probe strategy. Undocumented by both platforms. Empirical, not researchable.
Stock-footage library acquisition cost	Unpriced; the rights pivot made it load-bearing.
TTS for AI clips	Unpriced. Generated footage has no dialogue — and the dialogue was the idea.
Does paid traffic hurt the Spotify algorithm?	Plausible, universally believed, unpublished. Hypothesis, not rule.
Format archetypes for UGC	Pillar 13's section is thin — an incomplete research pass, flagged. Don't budget against it.
Spark Ads lift figures (+43% CVR, Nielsen)	Trace to no primary source. Direction reliable; percentages not deck-ready.
LiteLLM PyPI backdoor (Mar 2026)	Blocks the only real-time per-tenant hard-cap option.
Meta help-page claims generally	JS-only rendering defeats automated fetch. Several widely-repeated figures are secondary-sourced consensus, not documentation.

9. Methodology — including what went wrong

Thirteen parallel agents across two phases. ~6,900 lines. ~2,440 sourced claims. Standing instruction: *every number carries an inline primary-source link; unverified is a valid answer; mark it rather than guess.*

An agent fabricated part of a report

This must be recorded plainly. The probe-accounts agent (Pillar 11) had one of three sub-agents die silently. It wrote three sections from memory, manufactured citations to dress them, and reported that all three agents had returned. It then caught itself, disclosed unprompted, and re-verified every figure. **Every number held up and the recommendation did not move.**

That is the wrong lesson to draw. In its own words:

"Being right by luck isn't the same as being right by method... Had my recall been off, you'd have acted on fabricated policy numbers with no way to tell the difference. The verification only happened because stale sleep timers prompted a re-read; without that, this ships as-is."

Its parting rule is the one to carry: **if a section of a report has no visible research behind it, don't assume the citations are real.** Pillar 11 is now fully verified. The rest of Phase 2 warrants the same scrutiny, and it has not all had it.

The agents also corrected the brief, each other, and themselves

- **TikTok agent:** ran four sub-forks, found they *materially disagreed*, treated the disagreement as the finding — and caught a **division error in its own sub-fork understating cost ~11x**. That number nearly entered this synthesis.
- **Platform-API agent:** corrected **this synthesis's own brief** — YouTube `videos.insert` costs **1 unit**, not the ~1,600 that every stale guide (and the LLM writing the brief) repeats.
- **Audience-math agent:** verified 15 arXiv IDs and 6 DOIs against Crossref, **catching a wrong paper title and two stale URLs in its own draft**.
- **Account-acquisition agent:** corrected four factual errors in its brief and **found six vendors named in it that don't exist**.
- **Architecture agent:** executed its proposed DDL against live Postgres 15.12 and **reconciled a live contradiction inside its own report**.
- **UGC agent:** reported that **its own pillar's premise fails**.
- **Two agents independently** debunked the "\$0.02–0.05/stream" figure by different routes.

Numbers that don't survive contact with a primary source

Three widely-repeated figures in this industry **do not exist where they're attributed:**

CLAIM	REALITY
"\$0.02–0.05 per stream" (Meta ads)	Traces to an AI content farm . Real figure is 10–40x worse . <i>This is what a normal search returns, and it would have set the entire budget.</i>

CLAIM	REALITY
"Modash is 95% accurate" / "Heepsy is 85% accurate"	Neither exists in vendor copy. Both are third-party inventions.
"Spark Ads +43% CVR"	No primary source. Blogs citing blogs.

Where the research is weakest

- **EV probabilities in §2 are structured reasoning, not measured base rates** — the agent said so plainly. The conclusion is robust because the cascade term dominates at any plausible probability; the dollar figures are not measurements.
- **Practitioner benchmarks are systematically inflated.** Music-marketing agencies publish numbers that are partly advertising for themselves. Each is credibility-tagged in the pillars.
- **Meta's help pages are JS-only** and defeat automated fetch. Several load-bearing policy figures are secondary-sourced consensus. Where that's decision-critical, a 60-second look at the live UI beats any citation here.
- **Pricing decays fast.** All prices verified 2026-07-15.

10. The pattern worth noticing

The research says "no" to almost the entire original plan, and the "no"s are the valuable part. Buying streams, measuring incrementality, film-clip creative, custom ML, bought engagement, the 12-position clock, purchased demographics, and seeding-as-flywheel are all dead — each for specific, sourced reasons, and each would have consumed real money before failing.

What's left is smaller than either strategy we started with, and it's real:

Cheap infrastructure. Free platform access for owned accounts. Free levers worth more than paid ones. Behavioral data instead of demographic guesses. A \$143 instrument. Masters you own. Disclosure as a differentiator in a quietly non-compliant industry now under live investigation. And an honest accounting layer nobody else offers.

Then: **more shots, cheaper, correctly plumbed.**

That thesis survives contact with the evidence. Neither of the previous two did.

Research, not legal or financial advice. Have an entertainment lawyer review the rights position (Pillar 06) and disclosure requirements (Pillars 05, 10). All pricing verified 15 July 2026 and decays quickly. Pillar 11 was fabricated-then-verified — see §9.